

Common Python/Django Backend Interview Questions

Core Python

- What are `*args` and `**kwargs`?

Answer: In Python, `*args` collects extra positional arguments as a tuple, and `**kwargs` collects extra keyword arguments as a dict. Use them in function definitions to accept a flexible number of inputs. For example: `def func(a, b, *args, **kwargs): ...` lets callers pass additional positional or named arguments. Internally you can iterate over `args` or access `kwargs['key']` ¹.

Memory Aid: Think of `*args` as an extra “argument bag” (like stuffing all extra values into a single list) and `**kwargs` as a “key-value map” bag.

- Explain shallow copy vs deep copy.

Answer: A **shallow copy** (e.g. `new_list = old_list.copy()`) creates a new container but **references** the same nested objects. A **deep copy** (using `import copy; deep = copy.deepcopy(old)`) duplicates all levels, producing fully independent clones of nested items. Shallow is faster but modifying a nested object affects both copies; deep copy isolates them.

Memory Aid: Shallow copy = “clone skeleton” (same bones), deep copy = “clone with flesh” (everything duplicated).

- What’s the difference between a list and a tuple?

Answer: A *list* is mutable (you can add/remove/change elements) and is written with `[...]`. A *tuple* is immutable (fixed after creation) and is written with `(...)`. Use tuples for fixed collections (and they can be dict keys), and lists for dynamic collections. Tuples are slightly faster for iteration.

Memory Aid: Remember **MU**Table **LIST**, **IMM**utable **TU**ple (LIST you can “liSt”en and change, TUple you **turn** and it’s fixed).

- What is a decorator in Python and how do you use it?

Answer: A decorator is a function that wraps another function to modify its behavior without changing its code. You apply it with the `@decorator_name` syntax above a function definition. Under the hood, the decorator receives the original function as an argument and returns a new function. Common uses include logging, access control, or timing functions. For example, `@login_required` wraps a view to enforce authentication ¹.

Memory Aid: Think of a **gift wrap**: the decorator “wraps” a function with extra features. (Mnemonic: Decorator = **Don**ate functionality on top.)

- How do you handle exceptions in Python?

Answer: Use `try/except` blocks around code that might fail. Catch specific exceptions to handle known errors, e.g. `except ValueError:`. You can use an `else:` block for code that runs if no

exception occurs, and `finally:` for cleanup tasks (runs whether or not an exception happened). Avoid bare `except:` so you don't hide unexpected bugs. You can also raise exceptions with `raise`.

Memory Aid: Think “T-E-F”: **Try, Except, Finally** – like a safety net under trampoline: Try the code, Except handles tumbles, Finally lets you clean up afterwards.

- **What is a class method vs a static method?**

Answer: A `@classmethod` receives the class (`cls`) as its first argument. It can access or modify class state. A `@staticmethod` takes no special first argument – it neither receives `self` nor `cls` – and cannot access class/instance state. You use class methods for factory methods or operations that affect the class, and static methods for utility functions logically grouped in the class. For example, `@classmethod def from_csv(cls, file): ...` vs `@staticmethod def helpers(x): ...`.

Memory Aid: **Class method = “Class” first, Static = “Stand-alone.”** (Class methods get `cls`; static methods get nothing special.)

- **What is the Global Interpreter Lock (GIL)?**

Answer: In CPython, the Global Interpreter Lock (GIL) ensures only one thread executes Python bytecode at a time, even on multi-core machines. This means Python threads cannot execute CPU-bound code in parallel. I/O-bound tasks can still benefit from threading. For true CPU concurrency, you'd use multiprocessing or an asynchronous approach.

Memory Aid: Think of the GIL as a **traffic light** for threads: only one car (thread) can go at a time on the Python lane.

Django Framework Concepts

- **What is Django and what is its MVT architecture?**

Answer: Django is a full-stack Python web framework for rapid development of dynamic web applications ². It uses the Model-View-Template (MVT) pattern. *Models* define database schema (data layer), *Views* process requests and select data, and *Templates* render HTML. (Django's “View” is the Python code handling logic, and its “Template” is the HTML with placeholders). This separates data, business logic, and presentation ³.

Memory Aid: MVT is like MVC but with T for Template. **Model** = database/data, **View** = logic layer, **Template** = UI. (Think: Models do the “heavy lifting,” Views hand it to Templates for display.)

- **What is the difference between a Django project and an app?**

Answer: A *project* is the entire Django web application (the whole codebase), while an *app* is a modular component of that project with a specific purpose (e.g. a blog app, a user-auth app). A project can contain many apps, and apps are self-contained so they can be reused in other projects ⁴.

Memory Aid: Think of a **project** as a house, and **apps** as the individual rooms (kitchen, living room) each with a role.

- **How do you query the database using Django's ORM?**

Answer: Django's ORM exposes model classes with a manager (usually `objects`) to perform queries. For example, `MyModel.objects.all()` gets all rows;

`MyModel.objects.filter(field=value)` filters rows; `MyModel.objects.get(id=5)` retrieves one object (raises if not found); `MyModel.objects.create(...)` makes a new object and saves it. The ORM auto-constructs SQL and returns Python objects. You can chain filters and use methods like `.exclude()`, `.order_by()`, `.values()`, etc.

Memory Aid: Remember **CRUD** with the ORM: Create with `.create()`, Read with `.all()/.get()`, Update by changing fields and `.save()`, Delete with `.delete()`. (And DRY: use filters, selects.)

- **What is `select_related` and `prefetch_related`?**

Answer: These QuerySet methods optimize database access for related objects. `select_related('rel')` performs a SQL JOIN to fetch related objects in one query (useful for foreign-key “forward” relations). `prefetch_related('rel')` fetches related objects in a separate query and joins them in Python (useful for many-to-many or reverse relations). Both avoid the N+1 query problem by reducing the number of database hits.

Memory Aid: **Select = JOIN friend, Prefetch = separate fetch of group.** Think “select” for one big query, “prefetch” for two queries but still faster than many.

- **What are Django migrations?**

Answer: Migrations are Django’s way to propagate model (schema) changes to the database. When you change `models.py` (add fields, tables, etc.), you run `manage.py makemigrations` to auto-generate a migration file. Then `manage.py migrate` applies those schema changes to the database. Django treats migrations as **version control for your database schema** ⁵, ensuring everyone’s DB stays in sync.

Memory Aid: Think of migrations as **git commits for your database schema**: you “commit” model changes (makemigrations) and “push” to the DB (migrate).

- **What are Django middleware?**

Answer: Middleware are hooks in the request/response processing pipeline. Each middleware class can inspect or modify requests/responses. Common built-ins include security checks (`SecurityMiddleware`), session handling, CSRF protection (`CsrfViewMiddleware`), and authentication (`AuthenticationMiddleware`) ⁶. You configure them in `settings.py`; they run in order on every request.

Memory Aid: Think of middleware as **onion layers** around your view logic – each layer (middleware) can add processing before or after your views run.

- **What is the purpose of `urls.py` and `views.py`?**

Answer: `urls.py` maps URL patterns to view functions or classes. It defines routes, e.g. `path('books/', views.book_list)`. `views.py` contains the view code that handles a request and returns a response (often rendering a template). The URL router dispatches requests to the appropriate view based on the URL patterns.

Memory Aid: **URLs = map, Views = town guides.** The URL config maps a URL to a view that guides how to handle the request.

- **What is Django’s template system?**

Answer: Django templates are HTML files with placeholder tags (`{{ var }}`) and logic tags (`{% if %}`, `{% for %}`, etc.). They auto-escape variables for XSS protection. Templates allow

extending base layouts (`{% extends 'base.html' %}`) and including reusable fragments (`{% include %}`). They separate presentation from logic.

Memory Aid: Think of templates as **fillable forms**: static HTML with blanks (`{{ }}`) where Django inserts data.

- **What is Django REST Framework (DRF)?**

Answer: DRF is a toolkit for building Web APIs on top of Django. It provides serializers (for converting models to JSON), viewsets, routers, and authentication for RESTful endpoints. It greatly reduces boilerplate when creating JSON APIs. For example, DRF lets you declare a `ModelViewSet` and automatically get list/create/update/delete endpoints for a model.

Memory Aid: **DRF = Django's API factory.** (Like Django but geared for REST: think "DJANGO-Restified.") It's a "batteries-included" solution for APIs ⁷.

Database Handling (PostgreSQL, ORM, Migrations)

- **How do you configure and optimize a PostgreSQL database in Django?**

Answer: In `settings.py`, set `ENGINE: 'django.db.backends.postgresql'` and provide `NAME`, `USER`, `PASSWORD`, `HOST`, `PORT`. For optimization, ensure you use indexes on frequently filtered fields (`models.Index` or `db_index=True`), use `select_related` / `prefetch_related` to cut queries, and profile slow queries. You can use PostgreSQL-specific features (e.g. `JSONField` for JSON data). For large loads, use connection pooling (e.g. `CONN_MAX_AGE`) or external poolers.

Memory Aid: **"Index for speed, filter with care."** Always index columns you filter on, and avoid unnecessary queries.

- **What are migrations and how do you use them?**

Answer: (Combined here with optimization above.)

- **How do you prevent SQL injection in Django?**

Answer: Django's ORM automatically parameterizes queries, preventing injection. Avoid writing raw SQL. If you must, use `cursor.execute(sql, [params])` with parameters, never string-format user input. Always validate or sanitize user inputs (e.g. with Django Forms) before using them in queries or context.

Memory Aid: Think **ORM = bulletproof vest**. Using the ORM is your first defense against SQL injection.

- **How to perform database migrations in production without downtime?**

Answer: For simple changes (like adding non-null columns with a default or adding new tables), run migrations normally. For risky changes (e.g. dropping columns), use a rolling strategy: add a new field first (with `null=True`), deploy code that uses it, migrate data in batches, then later remove the old field. This two-phase approach avoids locking tables. Use `--noinput` and `--database` flags as needed and consider offline migrations for very large tables.

Memory Aid: **Two-phase migration:** Add first, deploy new code, then remove later (think: *add then drop*).

Security Best Practices

- **How does Django protect against Cross-Site Request Forgery (CSRF)?**

Answer: Django has built-in CSRF protection via middleware. It issues a unique token per user session and requires that each HTML form includes `{% csrf_token %}`. On form submission, Django checks the submitted token against the one in the user's cookies. If they match, the POST is allowed; otherwise it's rejected ⁸. This prevents attackers from forging form submissions.

Memory Aid: Think of the CSRF token as a **secret handshake code** in each form.

- **How does Django mitigate XSS (Cross-Site Scripting)?**

Answer: By default, Django templates auto-escape HTML special characters (`<`, `>`, `&`, etc.), turning user content harmless. You should not disable autoescape or use `mark_safe` unless you trust the content. Also, always validate/sanitize user input on forms ⁹. If you need to output raw HTML, explicitly mark it safe.

Memory Aid: Auto-escaping is like **replacing bullets with blanks** in user text so scripts can't execute.

- **How does Django handle SQL injection?**

Answer: When using the ORM or Django's query utilities, SQL queries are automatically parameterized, neutralizing malicious input. Only use raw SQL with parameter binding (`cursor.execute(sql, [params])`). Never concatenate user input into SQL strings.

Memory Aid: ORM = **parameterized queries** by default (like using `"?"` placeholders safely).

- **What about Clickjacking and HTTPS?**

Answer: Enable `X-Frame-Options` header (via `django.middleware.clickjacking.XFrameOptionsMiddleware`, default is `SAMEORIGIN`) to prevent your site from being framed. Always use HTTPS in production; set `SECURE_SSL_REDIRECT=True` and `SESSION_COOKIE_SECURE=True` / `CSRF_COOKIE_SECURE=True` so cookies only send over HTTPS ¹⁰. Also use HSTS (`SECURE_HSTS_SECONDS`) for strict transport security ¹¹.

Memory Aid: **"No frames, all secure."** Frame options and HSTS headers lock down your site's security footprint.

- **How should you manage secret keys and passwords?**

Answer: Never hard-code secrets in code. Use environment variables or a separate config (`python-decouple`, `django-environ`, etc.) to store `SECRET_KEY`, DB passwords, API keys, etc. Use Django's `make_password` and `check_password` for hashing user passwords (the default `User` model does this). Do not log sensitive information.

Memory Aid: Treat secrets like **nuclear codes**: store them off-limits and only bring them out when needed.

Testing in Django

- **How do you write unit tests in Django?**

Answer: Django uses Python's `unittest`. Write tests in `tests.py` (or a `tests/` directory) inside each app. Subclass `django.test.TestCase` which provides a test database. Use `self.client` to simulate requests (`self.client.get('/url/')`) and `assert` methods to

check responses and data. Test models, views, and any business logic. Run them with `manage.py test`.

Memory Aid: **TestCase = Test playground**. Each TestCase has its own test DB and client to play in.

- **What is pytest and how does it work with Django?**

Answer: `pytest` is an alternative testing framework. With the `pytest-django` plugin, you can write tests (even as plain functions) that access Django's ORM and settings. It offers a more concise syntax and powerful fixtures. To use it, install `pytest-django` and configure `pytest.ini`. You can then run `pytest` instead of `manage.py test`.

Memory Aid: **pytest = "pijavy test"**: simpler and more powerful. (Mnemonic: "Pytest says just write tests.")

- **What is the difference between unit tests and integration tests?**

Answer: *Unit tests* isolate a small piece (e.g. a single function or method) and mock external interactions. *Integration tests* exercise multiple components together (e.g. a view hitting the database). In Django, TestCase often mixes both, but true unit tests might patch out the DB and test a single model method. Integration tests would use the real ORM and full stack.

Memory Aid: **Unit = microscope, Integration = wide lens**. Unit tests zoom in, integration tests zoom out.

- **How do you test a Django view?**

Answer: Use `self.client` in a `TestCase`. For example, `response = self.client.get('/my-url/')`. Check `response.status_code`, `response.context`, or HTML content. For POST views, `self.client.post('/url/', data)`. You can also test class-based views by calling them directly if needed.

Memory Aid: Think of `self.client` as a **dummy browser** clicking around your site.

- **How do you measure test coverage?**

Answer: Use the `coverage` tool. Run tests with `coverage run manage.py test`, then `coverage report` to see coverage percentages. A high coverage (e.g. >80%) is often aimed for, but ensure meaningful tests, not just lines.

Memory Aid: Coverage is like a **test meter** – it shows how much of your code your tests "touch."

Deployment (Gunicorn, Nginx, Docker, etc.)

- **How do you deploy a Django application?**

Answer: In production, run Django behind a WSGI/ASGI server like Gunicorn or uWSGI. Commonly you put Nginx (or Apache) in front as a reverse proxy to serve static files and forward requests to Gunicorn. Use virtual environments for dependencies. Apply `collectstatic` to gather static files for Nginx. Configure the server for the correct `DJANGO_SETTINGS_MODULE` and use environment variables for secret configs.

Memory Aid: **Browser → Nginx → Gunicorn → Django**. Think of Nginx as the front door, Gunicorn as the gatekeeper, and Django as the butler handling requests.

- **What is Gunicorn and why use it?**

Answer: Gunicorn ("Green Unicorn") is a Python WSGI HTTP server. It runs multiple worker processes

that handle Django's WSGI app. It's lightweight, easy to configure, and widely used for Django deployment. It handles concurrent requests by forking worker processes.

Memory Aid: **Gunicorn = WSGI soldier**. It takes incoming HTTP and feeds it to your Django application.

- **Why use Nginx with Django?**

Answer: Nginx is a high-performance web server/reverse proxy. It efficiently serves static (CSS/JS/images) content, handles SSL termination, and proxies dynamic requests to Gunicorn. This offloads work from Gunicorn and improves performance and security.

Memory Aid: **Nginx = static champion**. Think of it as a traffic cop: it serves files directly and sends the rest to Django.

- **How would you containerize a Django app with Docker?**

Answer: Write a `Dockerfile` using an official Python base image. Copy your code, install dependencies (via `requirements.txt`), collect static files, and specify the Gunicorn command to run on container start. Use Docker Compose to also run your DB (Postgres) and cache (Redis) containers. Ensure to use environment variables for secrets and settings.

Memory Aid: **Docker = box it up**. Your app, its libs, and configs all packed into a container (like a portable shipping box).

- **What is ASGI and when would you use it?**

Answer: ASGI is the Asynchronous Server Gateway Interface, the successor to WSGI for async applications. Django 3+ supports ASGI, allowing async views and websockets (e.g. via Django Channels). You'd use an ASGI server like Daphne or Uvicorn if your app needs to handle asynchronous communication or WebSocket connections.

Memory Aid: **ASGI = Async WSGI**. If your app needs chat or long-lived connections, think ASGI.

Performance Scaling and Caching

- **How can you improve Django's performance?**

Answer: Several ways: enable **query caching** with `select_related` / `prefetch_related`; use Django's **cache framework** to store expensive page/template results (per-view or fragment cache)¹²; use Redis or Memcached as backend for caching and sessions; offload long tasks to background workers (Celery with Redis/RabbitMQ); and use database indexes. Profiling with tools like `django-silk` or just `EXPLAIN` on queries helps find bottlenecks.

Memory Aid: **Cache, Index, Scale**. (Think: fridge to cache pages, indexes to speed queries, Celery workers to offload tasks.)

- **What caching strategies does Django provide?**

Answer: Django has multiple caching levels¹³: *per-site cache* (cache all views), *per-view cache* (cache individual view outputs), *template fragment cache* (cache parts of templates), and *low-level cache API* (manually cache objects). You configure a cache backend (like Redis/Memcached) in `settings.CACHES`.

Memory Aid: Cache *all*, *view*, *piece*, or *manual* – like caching an entire photo, a single frame, a piece of it, or any object you want.

- **What is Celery and how does it relate to Django?**

Answer: Celery is an asynchronous task queue. In Django, use Celery to run time-consuming tasks (sending emails, processing files, etc.) in the background instead of on a web request. You configure a **broker** (often Redis or RabbitMQ) that holds the task queue. You then call tasks with `delay()` or `apply_async()`. This helps scale by not blocking the request/response cycle.

Memory Aid: **Celery = busy bees.** Bees (workers) run chores in the background so the web app doesn't slow down.

- **How do you use Redis with Django?**

Answer: Redis can serve as a cache or as a Celery broker. In `CACHES` settings, set a Redis cache backend (`django-redis`). For Celery, configure `CELERY_BROKER_URL = 'redis://...'`. You can also use Redis for Django sessions or rate-limiting. Redis is an in-memory store, so it's much faster than hitting the DB for repeated data.

Memory Aid: **Redis = speed memory.** Like a high-speed interim storage (RAM) for frequently used data.

- **What are some ways to scale a Django application?**

Answer: Scale vertically (bigger server) or horizontally (multiple app servers behind a load balancer). Use a separate DB server or a managed DB, and separate Redis/Cache servers. Serve static files via CDN. Use Django's cache for heavyweight computations. Break tasks into microservices if needed. Also optimize code (remove N+1 queries, use pagination).

Memory Aid: **"More boxes."** More app instances + better caching = handling more users. (Think: add more workers or bigger machines.)

Best Practices and Design Patterns in Django

- **What is "fat models, thin views" principle?**

Answer: Keep business logic and data manipulation in model methods (or separate service classes), not in views. Views should orchestrate (call model methods, render templates) but not contain complex logic. This makes code reusable and easier to test.

Memory Aid: **Fat Model, Skinny View.** Pack the logic into models; views just "serve" it.

- **How do you handle settings for different environments (development, staging, production)?**

Answer: Don't hard-code settings. Use environment variables or separate settings modules. For example, use `python-decouple` or `django-environ` to read a `.env` file. Alternatively, have `settings/` package with `dev.py`, `prod.py`, and activate the right one with `DJANGO_SETTINGS_MODULE`. Keep `SECRET_KEY`, debug flags, and DB credentials environment-specific.

Memory Aid: **"Two settings are better than one."** Or more – one for dev, one for prod. (Mnemonic: *DJP* – Dev/Job/Prod separate.)

- **What are some common Django design patterns?**

Answer: Common patterns include *generic views* (class-based generic views like `ListView`, `CreateView` to avoid boilerplate), *signals* for decoupled events (e.g. `post_save` to trigger actions), *middleware* for cross-cutting concerns, and *context processors* for injecting data into all templates. Also using `Manager` classes for complex queries, and `ModelForm` for validation.

Memory Aid: Think **DRY** (Don't Repeat Yourself): use Django's built-in abstractions (generic CBVs, forms, signals) instead of re-coding logic.

- **What is a Django signal and when would you use it?**

Answer: Signals let you hook into model events. For example, `post_save` runs after a model's `save()`, `post_delete` after deletion. Use them to trigger related actions (e.g. create a user profile after a User is created) ¹⁴. However, they can make flow harder to trace, so use judiciously.

Memory Aid: **Signal = event listener**. Like an alarm that triggers extra jobs when certain model events happen.

- **How do you handle static and media files?**

Answer: Static files (CSS/JS/images for the site itself) go into `STATIC_ROOT` (collected via `collectstatic`) and are served by the web server. User-uploaded files (media) go into `MEDIA_ROOT` and served via `MEDIA_URL` (in dev, Django can serve them; in prod, usually Nginx or a cloud storage bucket serves media). In settings:

```
STATIC_URL = '/static/'; STATIC_ROOT = '/var/www/static/'  
MEDIA_URL = '/media/'; MEDIA_ROOT = '/var/www/media/'
```

Memory Aid: **Static = site furniture; Media = user furniture**. Static is built-in; media is made by users.

- **What best practices do you follow in Django development?**

Answer: Keep code DRY, follow PEP 8 for style. Use a virtual environment, pin dependencies in `requirements.txt`. Use the built-in `User` model or a clear custom one. Secure all forms and views. Write tests. Use logging and handle exceptions. Keep settings out of version control. Use transactions (`atomic`) for multi-step DB operations. Update Django and libraries for security.

Memory Aid: **KISS and DRY**: Keep It Simple, Stupid; Don't Repeat Yourself. (Mnemonic: *LOGIC* – Lint/check code style, Optimize queries, Guard against attacks, Include tests, Clean architecture.)

Sources: Authoritative Django docs and community Q&A were used to ensure accuracy and current best practices ² ³ ⁵ ⁸ ¹² ¹ ¹⁵.

¹ Primer on Python Decorators – Real Python

<https://realpython.com/primer-on-python-decorators/>

² ³ ⁴ ⁶ ⁷ ⁸ ¹² ¹³ ¹⁴ Top 50 Django Interview Questions and Answers | GeeksforGeeks

<https://www.geeksforgeeks.org/django-interview-questions/>

⁵ Migrations | Django documentation | Django

<https://docs.djangoproject.com/en/5.2/topics/migrations/>

⁹ ¹⁰ ¹¹ Security in Django | Django documentation | Django

<https://docs.djangoproject.com/en/5.2/topics/security/>

¹⁵ Class method vs Static method in Python | GeeksforGeeks

<https://www.geeksforgeeks.org/class-method-vs-static-method-python/>